



```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```


```
data = pd.read_csv('/content/train.csv')
data.head()
```



	<b>Id</b>	<b>MSSubClass</b>	<b>MSZoning</b>	<b>LotFrontage</b>	<b>LotArea</b>	<b>Street</b>	<b>Alley</b>	<b>LotShape</b>	<b>LandContour</b>
<b>0</b>	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl
<b>1</b>	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl
<b>2</b>	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl
<b>3</b>	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl
<b>4</b>	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl

5 rows × 81 columns





data.columns

```
Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
      'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
      'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
      'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
      'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
      'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
      'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
      'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
      'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
      'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
      'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
      'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
      'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
      'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
      'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
      'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
      'SaleCondition', 'SalePrice'],
      dtype='object')
```

data.shape

```
(1460, 81)
```

```
df = data[['OverallQual', 'GrLivArea', 'TotalBsmtSF', '1stFlrSF',
          'FullBath', 'YearBuilt', 'YearRemodAdd', 'GarageCars', 'GarageArea',
          'Fireplaces', 'LotArea', 'ExterQual', 'KitchenQual', 'BsmtQual', 'SalePrice']]
df.head() #These Columns are helpful for house price prediction
```

	<b>OverallQual</b>	<b>GrLivArea</b>	<b>TotalBsmtSF</b>	<b>1stFlrSF</b>	<b>FullBath</b>	<b>YearBuilt</b>	<b>YearRemodAdd</b>	<b>Gar</b>
<b>0</b>	7	1710	856	856	2	2003	2003	
<b>1</b>	6	1262	1262	1262	2	1976	1976	
<b>2</b>	7	1786	920	920	2	2001	2002	
<b>3</b>	7	1717	756	961	1	1915	1970	
<b>4</b>	8	2198	1145	1145	2	2000	2000	

Next steps:

Generate code with df

 View recommended plots

df.shape

```
(1460, 15)
```

df.columns

```
Index(['OverallQual', 'GrLivArea', 'TotalBsmtSF', '1stFlrSF', 'FullBath',
      'YearBuilt', 'YearRemodAdd', 'GarageCars', 'GarageArea', 'Fireplaces',
      'LotArea', 'ExterQual', 'KitchenQual', 'BsmtQual', 'SalePrice'],
      dtype='object')
```

df.isnull().sum()

```
OverallQual    0
GrLivArea      0
TotalBsmtSF    0
1stFlrSF       0
FullBath       0
YearBuilt      0
YearRemodAdd   0
GarageCars     0
GarageArea     0
Fireplaces     0
LotArea        0
ExterQual      0
KitchenQual    0
BsmtQual       37
SalePrice      0
dtype: int64
```

```
import warnings
warnings.filterwarnings('ignore')
df.dropna(inplace=True)
```

```
df.isnull().sum() #Now there is no missing values.
```

```
OverallQual    0
GrLivArea      0
TotalBsmtSF    0
1stFlrSF       0
FullBath       0
YearBuilt      0
YearRemodAdd   0
GarageCars     0
GarageArea     0
Fireplaces     0
LotArea        0
ExterQual      0
KitchenQual    0
BsmtQual       0
SalePrice     0
dtype: int64
```

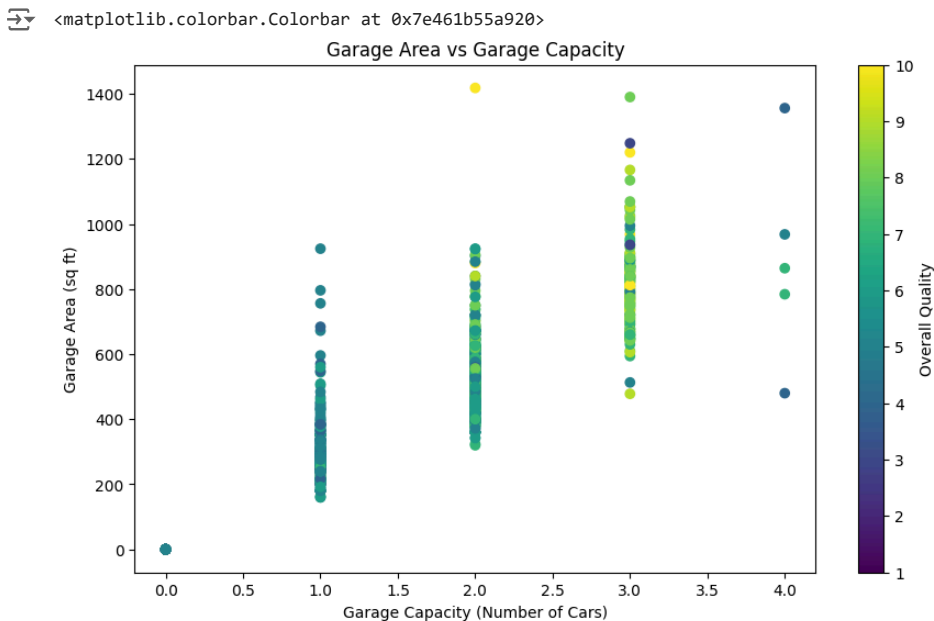
```
df.duplicated().sum() #No Duplicates are there
```

```
0
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1423 entries, 0 to 1459
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   OverallQual     1423 non-null   int64
1   GrLivArea       1423 non-null   int64
2   TotalBsmtSF     1423 non-null   int64
3   1stFlrSF        1423 non-null   int64
4   FullBath        1423 non-null   int64
5   YearBuilt       1423 non-null   int64
6   YearRemodAdd    1423 non-null   int64
7   GarageCars      1423 non-null   int64
8   GarageArea      1423 non-null   int64
9   Fireplaces      1423 non-null   int64
10  LotArea         1423 non-null   int64
11  ExterQual       1423 non-null   object
12  KitchenQual     1423 non-null   object
13  BsmtQual        1423 non-null   object
14  SalePrice       1423 non-null   int64
dtypes: int64(12), object(3)
memory usage: 177.9+ KB
```

```
plt.figure(figsize=(10, 6))
plt.scatter(df['GarageCars'], df['GarageArea'], c=df['OverallQual'])
plt.xlabel('Garage Capacity (Number of Cars)')
plt.ylabel('Garage Area (sq ft)')
plt.title('Garage Area vs Garage Capacity')
plt.colorbar(label='Overall Quality')
```



Observations:


1. For a garage with the capacity for 2 cars and an excellent area (in square feet), the quality is rated as very good, 10/10.

2. For a garage with the capacity for 3 cars and a good area, the quality is rated as good, ranging from 8.5 to 9 out of 10.

```
numericals = df.select_dtypes(include=['float64', 'int64'])
q1 = numericals.quantile(0.25)
q3 = numericals.quantile(0.75)
IQR = q3-q1
lower = q1 - 1.5*(IQR)
higher = q3 + 1.5*(IQR)
outliers = df[((numericals < lower) | (numericals > higher)).any(axis=1)]
print("Total Outliers in this dataset : ",outliers.shape[0])
```

Total Outliers in this dataset : 160

```
plt.figure(figsize=(15,7))
sns.boxplot(df)
```



```
numericals = df.select_dtypes(include=['float64', 'int64'])
q1 = numericals.quantile(0.25)
q3 = numericals.quantile(0.75)
IQR = q3-q1
lower = q1 - 1.5*(IQR)
higher = q3 + 1.5*(IQR)
df1 = df[~((numericals < lower) | (numericals > higher)).any(axis=1)]
df1.shape
```

(1263, 15)

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1263 entries, 0 to 1459
Data columns (total 15 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   OverallQual     1263 non-null  int64  
 1   GrLivArea       1263 non-null  int64  
 2   TotalBsmntSF    1263 non-null  int64  
 3   1stFlrSF        1263 non-null  int64  
 4   FullBath        1263 non-null  int64  
 5   YearBuilt       1263 non-null  int64  
 6   YearRemodAdd    1263 non-null  int64  
 7   GarageCars      1263 non-null  int64  
 8   GarageArea      1263 non-null  int64  
 9   Fireplaces      1263 non-null  int64  
10  LotArea         1263 non-null  int64  
11  ExterQual       1263 non-null  object  
12  KitchenQual     1263 non-null  object  
13  BsmtQual        1263 non-null  object  
14  SalePrice       1263 non-null  int64  
dtypes: int64(12), object(3)
memory usage: 157.9+ KB
```

```
df1['ExterQual'].unique()
```

array(['Gd', 'TA', 'Ex', 'Fa'], dtype=object)

```
df1['KitchenQual'].unique()
```

array(['Gd', 'TA', 'Fa', 'Ex'], dtype=object)

```
df1['BsmtQual'].unique()
```

array(['Gd', 'TA', 'Ex', 'Fa'], dtype=object)

Ex (Excellent)

Gd (Good)

TA (Typical/Average)

Fa (Fair)

```
from sklearn.preprocessing import OrdinalEncoder
order = ['Fa','TA','Gd','Ex'] #Lowest to highest
encoder = OrdinalEncoder(categories=[order,order,order])
df1[['ExterQual','KitchenQual','BsmtQual']]=encoder.fit_transform(df1[['ExterQual','KitchenQual','BsmtQual']])
df1.head()
```

	OverallQual	GrLivArea	TotalBsmtSF	1stFlrSF	FullBath	YearBuilt	YearRemodAdd	Gar
0	7	1710	856	856	2	2003	2003	
1	6	1262	1262	1262	2	1976	1976	
2	7	1786	920	920	2	2001	2002	
3	7	1717	756	961	1	1915	1970	
4	8	2198	1145	1145	2	2000	2000	

Next steps:

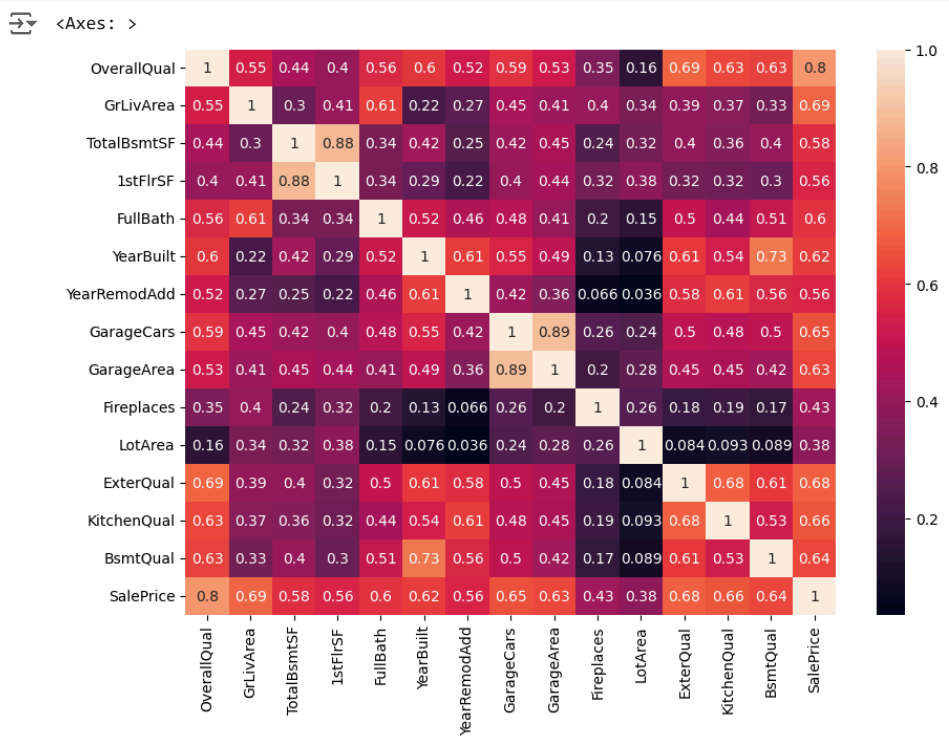
[Generate code with df1](#)

☒ [View recommended plots](#)

```
df1.describe()
```

	OverallQual	GrLivArea	TotalBsmtSF	1stFlrSF	FullBath	YearBuilt	YearR
count	1263.000000	1263.00000	1263.000000	1263.000000	1263.000000	1263.000000	1263.000000
mean	6.012668	1435.81631	1028.661916	1105.018211	1.520190	1971.153603	1984.018211
std	1.237232	418.77310	325.925913	314.878605	0.524544	29.748979	20.748979
min	2.000000	438.00000	105.000000	438.000000	0.000000	1885.000000	1950.000000
25%	5.000000	1114.00000	797.000000	864.000000	1.000000	1954.000000	1966.000000
50%	6.000000	1414.00000	971.000000	1053.000000	2.000000	1973.000000	1993.000000
75%	7.000000	1705.50000	1230.000000	1309.500000	2.000000	2000.000000	2003.000000
max	10.000000	2730.00000	2000.000000	2117.000000	3.000000	2009.000000	2010.000000

```
plt.figure(figsize=(10,7))
sns.heatmap(df1.corr(),annot=True)
```



Observation: Every column has a good correlation with the target variable, so we can't remove any columns...

```
details = ['OverallQual', 'GrLivArea', 'TotalBsmtSF', '1stFlrSF','YearBuilt', 'YearRemodAdd', 'GarageArea', 'LotArea', 'SalePrice']
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df1[details] = scaler.fit_transform(df1[details])
df1.head()
```



	OverallQual	GrLivArea	TotalBsmtSF	1stFlrSF	FullBath	YearBuilt	YearRemodAdd	Gar
0	0.625	0.554974	0.396306	0.248958	2	0.951613	0.883333	
1	0.500	0.359511	0.610554	0.490768	2	0.733871	0.433333	
2	0.625	0.588133	0.430079	0.287076	2	0.935484	0.866667	
3	0.625	0.558028	0.343536	0.311495	1	0.241935	0.333333	
4	0.750	0.767888	0.548813	0.421084	2	0.927419	0.833333	

Next steps:

[Generate code with df1](#)



[View recommended plots](#)

```
x=df1.iloc[:, :-1] #Independent
y=df1.iloc[:, -1] #Dependent
x
```



	OverallQual	GrLivArea	TotalBsmtSF	1stFlrSF	FullBath	YearBuilt	YearRemodAdd	
0	0.625	0.554974	0.396306	0.248958	2	0.951613	0.883333	
1	0.500	0.359511	0.610554	0.490768	2	0.733871	0.433333	
2	0.625	0.588133	0.430079	0.287076	2	0.935484	0.866667	
3	0.625	0.558028	0.343536	0.311495	1	0.241935	0.333333	
4	0.750	0.767888	0.548813	0.421084	2	0.927419	0.833333	
...	...	...	...	...	...	...	...	...
1455	0.500	0.527487	0.447493	0.306730	2	0.919355	0.833333	
1456	0.500	0.713351	0.758311	0.973794	2	0.750000	0.633333	
1457	0.625	0.829843	0.552507	0.446694	2	0.451613	0.933333	
1458	0.375	0.279232	0.513456	0.381179	1	0.524194	0.766667	
1459	0.375	0.356894	0.607388	0.487195	1	0.645161	0.250000	

1263 rows × 14 columns



Next steps:

[Generate code with x](#)



[View recommended plots](#)

y



```
0    0.569460
1    0.480892
2    0.618665
3    0.344760
4    0.705593
...
1455 0.459570
1456 0.574381
1457 0.759718
1458 0.351730
1459 0.369362
Name: SalePrice, Length: 1263, dtype: float64
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=0)
```

```
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(x_train,y_train)
```



LinearRegression

LinearRegression()

```
from sklearn.metrics import r2_score,mean_absolute_error
y_pred = lr.predict(x_test)
print(r2_score(y_test,y_pred))
print(mean_absolute_error(y_test,y_pred))
```



```
0.866458713200901
0.05133304229899392
```

```
from sklearn.ensemble import RandomForestRegressor
rfr = RandomForestRegressor()
rfr.fit(x_train,y_train)
```



RandomForestRegressor

RandomForestRegressor()

```
y_pred = rfr.predict(x_test)
print(r2_score(y_test,y_pred))
print(mean_absolute_error(y_test,y_pred))
```

0.8580312311610334  
0.04963907350747254

Start coding or [generate](#) with AI.

Conclusion :

Linear Regression : 86.6 %

Random Forest : 85.4%

By hypertuning using grid search, I found that the accuracy did not improve significantly. The highest accuracy achieved was 86.6% using linear regression model.

Start coding or [generate](#) with AI.

Testing Part:

```
new = pd.DataFrame([[7,1710,856,856,2,2003,2003,2,548,0,8450,2.0,2.0,2.0]],columns=x.columns)
new
```

	OverallQual	GrLivArea	TotalBsmtSF	1stFlrSF	FullBath	YearBuilt	YearRemodAdd	GarageCars	GarageArea	Fireplaces	LotArea	ExterQual
0	7	1710	856	856	2	2003	2003	2	548	0	8450	

```
output = lr.predict(new)
print("Sales Price of house is : ",output)
```

Sales Price of house is : [2070.92384172]

Start coding or [generate](#) with AI.